

CMPUT463/563
Probabilistic Graphical Models

Exact Inference: Variable Elimination

Lili Mou

Dept. Computing Science, University of Alberta

lmou@ualberta.ca

Outline

- Types of probabilistic queries
- Naïve calculation is expensive
- Dynamic programming on a chain
- Variable elimination in general
 - Evidence potentials falsify incompatible assignments
 - Eliminate a variable by fully connecting its neighbors
 - Sum-product and max-product are semirings
 - Directed and undirected graphs work in a similar way
- Efficiency of VE depends on the induced-width
 - Finding the best order is NP-hard
 - Intuition helps

Inference Questions

We group the variables of a sample into three sets X, Y, Z

- X : observed; Y : variables in question
- Z : to be marginalized out

Query types

- Probabilistic queries: $P(Y | X)$
 - Special case $Y = \emptyset$: probability of evidence $P(X)$
- Max a posteriori (MAP) inference: $\operatorname{argmax} P(Y | X)$
 - Most likely values for **some** variables
 - Special case
 - Most probable explanation (MPE): when $Z = \emptyset$
 - Most likely values for **all other** variables

Inference Questions - Applications

Query types


- Probabilistic queries: $P(Y | X)$
 - Special case $Y = \emptyset$: probability of evidence $P(X)$

Application: Outlier detection; Used in parameter learning

- Max a posteriori (MAP) inference: $\operatorname{argmax} P(Y | X)$
 - Most likely values for **some** variables
 - Special case
 - Most probable explanation (MPE):
 $\operatorname{argmax}_{Y,Z} P(Y, Z | X)$
 - Most likely values for **all other** variables

MAP and MPE may be different for Y

Applications: Image segmentation, POS tagging
Text generation from continuous latent space



Joint probability tells everything

- Probabilistic queries: $P(Y | X)$

$$P(Y | X) \propto_Y \sum_z P(Y, z | X)$$

- Max a posteriori (MAP) inference: $\operatorname{argmax} P(Y | X)$

$$\operatorname{argmax}_Y \sum_z P(Y, z | X)$$

- Most probable explanation (MPE): when $Z = \emptyset$


$$\operatorname{argmax}_Y P(Y | X)$$

Complexity of naïve computation?

Categorization of Inference Algorithms

- Exact inference
 - Variable elimination, message passing, junction tree
 - With dynamic programming, efficiency is better than enumeration, but still NP-hard (NP-complete or harder) for a general graph
 - Tree structures: linear
- Approximate inference
 - Asymptotically correct: Monte Carlo approaches
 - Deterministically wrong: variational inference


Why dynamic programming is possible?



- Consider the query $P(e)$


$$P(e) = \sum_d \sum_c \sum_b \sum_a P(a)P(b|a)P(c|b)P(d|c)P(e|d)$$

$$= \sum_d \sum_c \sum_b \left[\underbrace{\left(\sum_a P(a)P(b|a) \right)}_{P(b)} P(c|b)P(d|c)P(e|d) \right]$$



Variable A is eliminated

Why dynamic programming is possible?



- Consider the query $P(e)$


$$P(e) = \sum_d \sum_c \sum_b P(b)P(c|b)P(d|c)P(e|d)$$

$P(c)$


$$= \sum_d \sum_c P(c)P(d|c)P(e|d)$$

$P(d)$

$$= \sum_d P(d)P(e|d)$$



Such DP also works for argmax



- Consider the query $\max_{a,b,c,d} P(e)$ Subscripts of $m(\cdot)$ indicate they're different functions

$$\max_{a,b,c,d} P(e) = \max_d \max_c \max_b \max_a P(a)P(b|a)P(b)P(c|b)P(d|c)P(e|d)$$

a $m_B(b)$

$$= \max_d \max_c \max_b m(b)P(c|b)P(d|c)P(e|d)$$

b $m_C(c)$


$$= \max_d \max_c m(c)P(d|c)P(e|d)$$

c $m_D(d)$

$$= \max_d m(d)P(e|d)$$

d

This is dynamic programming (DP), not a greedy algorithm. For example, when computing $\max_c P(c)P(d|c)$, you get a function of d without a specific c . The choice of c depends on d , so back-pointers are needed.



Semiring

Compared with a ring, semiring
does not require inverse of addition

- Algebraic structure: (\oplus, \odot) on a set R is a semiring if
 - \oplus is associative and commutative with identity 0
 - $0 \oplus a = a \oplus 0 = a$
 - $(a \oplus b) \oplus c = a \oplus (b \oplus c)$
 - $a \oplus b = b \oplus a$
 - \odot is associative with zero 0 and identity 1
 - $0 \odot a = a \odot 0 = 0, 1 \odot a = a \odot 1 = a$
 - $(a \odot b) \odot c = a \odot (b \odot c)$
 - \odot distributive wrt \oplus
 - $$a \odot (b \oplus c) = (a \odot b) \oplus (a \odot c)$$
$$(b \oplus c) \odot a = (b \odot a) \oplus (c \odot a)$$

Semiring


- Both $(+, \cdot)$ and (\max, \cdot) are semirings
- The above DP algorithms only involve the operations defined in a semiring
- Thus, the two algorithms are almost identical.

$$P(e) = \sum_d \sum_c \sum_b \sum_a P(a)P(b|a)P(c|b)P(d|c)P(e|d)$$

$$\max_{a,b,c,d} P(e) = \max_d \max_c \max_b \max_a P(a)P(b|a)P(b)P(c|b)P(d|c)P(e|d)$$

To obtain argmax, we need back-pointers

Such DP also works for undirected graphs



$$\begin{aligned}
 P(e) &= \sum_d \sum_c \sum_b \sum_a \frac{1}{Z} \phi_1(a, b) \phi_2(b, c) \phi_3(c, d) \phi_4(d, e) \\
 &= \frac{1}{Z} \sum_d \sum_c \sum_b \sum_a \phi_1(a, b) \phi_2(b, c) \phi_3(c, d) \phi_4(d, e)
 \end{aligned}$$

$m_B(b)$
 $m_C(c)$
 $m_D(d)$
 $m_D(d)$

How about Z ?

- Sum-product: marginalize out E
- Max-product: can be ignored because Z is a constant

Beyond a chain structure

From now on, we consider undirected graphs

For BNs, a conditional probability $P(X | \text{Par}(X))$ can immediately be interpreted as $\phi(X, \text{Par}(X))$

Suppose our query is $P(Y | x)$ and we need to eliminate \mathbf{Z}

$$P(Y | x) \propto P(Y, x) = \sum_{\mathbf{z}} P(Y, \mathbf{z}, x)$$


$$\propto \sum_{\mathbf{z} \setminus \{z_0\}} \underbrace{\sum_{z_0} \prod_{\phi: z_0 \in \text{scope}(\phi)} \phi}_{\text{red bracket}} \prod_{\psi: z_0 \notin \text{scope}(\psi)} \psi$$

$$\phi \text{ with scope } \cup \text{scope}(\phi) \setminus \{z_0\}$$
$$\phi: z_0 \in \text{scope}(\phi)$$

Eliminate a Variable

To eliminate a variable Z_0 :

- Consider all factors involving Z_0
- Eliminate Z_0 by summing all the values of Z_0
- Obtain a factor φ with a scope of Z_0 ' neighbors (fully connects the neighbor)
- Put φ back to the graphical model




Handling Evidence

$P(Y|x) \propto P(Y, x)$, i.e., we only considers assignments compatible to evidence.

This can be handled by introducing an evidence potential

$$\delta_i = \begin{cases} 1, & \text{if } X_i = x_i \\ 0, & \text{if } X_i \neq x_i \end{cases}$$

Multiply every potential containing X_i with δ_i



X_i	...	ϕ
0	0	2
0	1	0.1
1	0	0.2
1	1	10

X_i	...	$\phi\delta_i$	
0	0	2	0
0	1	0.1	0
1	0	0.2	
1	1	10	

Variable Elimination

Input: Variables $Y, Z, X = x$, set of factors Φ

Elimination order (assuming Z_1, Z_2, \dots, Z_k wlog)

1. $\Phi \leftarrow \Phi$ with evidence potentials

2. For $i = 1, \dots, k$


2.1 $\Phi' \leftarrow \{\phi \in \Phi : Z_i \in \text{scope}(\phi)\}$

2.2. $\varphi \leftarrow \sum_{z_i} \prod_{\phi \in \Phi'} \phi$


2.3 $\Phi = \Phi \setminus \Phi' \cup \{\varphi\}$

Return $\prod_{\phi \in \Phi} \phi$ as a factor on Y given X


Some graphs don't have efficient elimination anyway




Certain graphs can be eliminated efficiently



Chordal



Example




$$\begin{aligned}
 P(e) &= \sum_d \sum_c \sum_b \sum_a \frac{1}{Z} \phi_1(a, b) \phi_2(b, c) \phi_3(c, d) \phi_4(d, e) \\
 &= \frac{1}{Z} \sum_d \sum_c \sum_b \sum_a \phi_1(a, b) \phi_2(b, c) \phi_3(c, d) \phi_4(d, e)
 \end{aligned}$$

$\underbrace{\hspace{10em}}_{m_B(b)}$
 $\underbrace{\hspace{10em}}_{m_C(c)}$
 $\underbrace{\hspace{10em}}_{m_D(d)}$
 $\underbrace{\hspace{10em}}_{m_D(d)}$

In fact, you may eliminate either direction in this example. However, the eliminator order matters a lot

Example



Query: $P(X_k | X_0)$

Eliminate X_0, \dots, X_{k-1}, Z : linear complexity

Eliminate Z first: exponential wrt k

The *induced-width* is a size of the largest scope during variable elimination given an induction order. The *tree-width* is the minimum induced width.

Unfortunately, $\text{TreeWidth} \leq M$ is NP-complete. Finding the smallest tree-width is NP-hard. Intuition helps.

Summary

- Types of probabilistic queries
- Naïve calculation is expensive
- Dynamic programming on a chain
- Variable elimination in general
 - Evidence potentials falsify incompatible assignments
 - Eliminate a variable by fully connecting its neighbors
 - Sum-product and max-product are semirings
 - Directed and undirected graphs work in a similar way
- Efficiency of VE depends on the induced-width
 - Finding the best order is NP-hard
 - Intuition helps